

Trillet.ai — Response & Operating Plan

My response to the audit review, backed with data where I could put a number behind it, plus a sorted 7-day / 30-day plan spanning technical, content revamp, main-page positioning, interlinking, new content, and marketplace profiles — with product-and-docs understanding built in before the content work.

Prepared — July 2026 (v3)

Sources — GSC (Web; 12-month + Sep–Nov 2025 & Mar–May 2026), robots.txt, sitemap.xml, page view-source, GA4

Access — read-only; scoped changes via dev PR workflow

VERIFIED from Trillet's own data/source **INFERRED** reasoned / experience-based **UNVERIFIED** needs access not yet granted

My response to the review

Addressed point by point. Where I could put a number behind an observation I have; where it's a judgment call from experience, I've said so. Full labels throughout, since crawling was blocked on the first pass. (A longer written version accompanies this document.)

1 · "No organic assets" vs. brand share.

The traffic exists — my point was proportional. Non-brand is **~3.5% of total clicks** **VERIFIED**, under 6% even with the export cap, and much of it top-of-funnel. So proportionally our organic presence is sparse and scattered — we rank widely, capture little non-brand demand.

2 · Seven hostnames.

Two things bundled. **Apex is the technical priority** (confirmed, low-risk). The subdomains (docs/security/certification/compliance) are a *qualitative* AI-visibility call — they're crawled as separate domains and hold the content AI crawlers value, so consolidating *could* help, but it's a choice on timing, not a mandate. → the "evaluate, then decide" column.

3 · Purging old URLs.

301 + a Search Console removal request. The 301 does the durable consolidation; the removal request just accelerates de-indexing so crawl reallocates to important pages sooner. Removal is temporary, so it's used alongside the 301, never instead.

4 · Crawl budget.

Fair — not "budget" in the ecommerce sense. It's a **crawl-efficiency / index-freshness** point: low-value URLs spread crawl thin and, in my experience, slow how fast important-page changes get indexed **INFERRED**.

5 · CTR decline — cohorts.

Ran the two windows as a matched cohort: **93% of the drop is composition** (292 new pages at 0.20% CTR, now 91% of impressions), while the stable cohort held CTR (3.98% → 3.74%) and position (5.32 → 5.23) **VERIFIED**. Not a site-wide snippet failure — a scoped opportunity.

6 · The ~4,000 clicks.

A quantification, not a forecast. Striking-distance non-brand pages (pos ~3–20) carry **155,009 impressions** **VERIFIED**; at 2–2.5% CTR that's ~3,100–3,875 clicks (~2,700–3,450 uplift). Caveat: current CTR there is ~0.4%, so 2–2.5% is the "if snippets reach normal" ceiling — and the base is capped, so it's a floor on the opportunity.

7 · "Answering service" terms.

From running this for Insighto.ai — I don't read it as generic. The intent maps onto "ai answering service"; the searcher wants the outcome and needs a reason to switch manual → AI. Data agrees near-term: the commercial long-tail (cheap/roofing/plumbing answering service) already ranks page 1 **VERIFIED**. Capture that now; build toward the broader term over a longer horizon.

8 · AEO.

No settled playbook exists, so I won't build the plan as if there is one. **Test-measure-scale**: run schema / answer-first / llms.txt experiments, watch AI-referral traffic and citations, scale what works. (FAQ rich results are gov/health-only now; HowTo is retired — value is machine-readability, which is why it's worth testing, not assuming.)

9 · Full funnel.

Depends on attribution, event config, and data volume. Non-brand traffic is thin and GA4 key events read 0 **VERIFIED** — so deep funnel/payment conclusions today would over-read noise. **Instrument now, analyze when volume justifies it.**

10–12 · Plan / labels / source of truth.

Plan attached with the full field set per action, sorted. Every conclusion labeled. Every pricing/product/compliance claim gated behind an approved source before publishing — which is why the plan sets aside time to understand the product and docs first, so content is grounded in what the platform actually does.

Part Two · 7-Day plan (sorted)

Sequenced by what unblocks what: understand the product, lock the source of truth, evaluate the technical fixes, stand up measurement, and map the content work. Read-only or staged except the scoped changes, which route through the dev PR workflow on access + sign-off.

7D·1 Product & docs deep-dive — understand before writing		Days 1–3
ASSET	docs.trillet.ai, the app, and feature set: honeypot detection, multi-agent orchestration (crews), website-scrape onboarding, MCP/integrations, on-premise Docker, compliance suite, human handoff.	
EVIDENCE	Content claims currently paraphrase marketing, not verified capability UNVERIFIED .	
PRODUCT MOTION	Study docs + product; produce a one-page "verified capabilities & differentiators" reference.	
MECHANISM	Content and positioning written against real capability, not restated copy — feeds the source-of-truth register.	
OWNER	[Content/SEO] with [Product]	
DEPENDENCY	Docs access (public) + 30 min with product.	
SUCCESS METRIC	Capabilities reference signed off by [Product]; differentiators mapped to target queries.	
WINDOW	Day 3.	
FAILURE CONDITION	If a claimed capability can't be verified, it's excluded from content until it can.	

7D·2 Source-of-truth register + claim owners		Days 1–5
ASSET	Register of pricing / compliance / product claims for published content.	
EVIDENCE	All claims trace to marketing surfaces UNVERIFIED as canonical .	
PRODUCT MOTION	Governance; each claim → approved owner; compliance → legal.	
MECHANISM	Protects the credibility of the pages we optimize.	
OWNER	[Product Marketing] + [Compliance/Legal]	
DEPENDENCY	Named owners; 7D·1 capabilities reference.	
SUCCESS METRIC	100% of to-be-published claims have an approved source.	
WINDOW	Day 5.	
FAILURE CONDITION	Unowned claims excluded from content.	

7D·3 Confirm the meta-truncation bug is site-wide

Days 1–2

ASSET	Meta-generation template; 8–10 sampled pages.
EVIDENCE	Description cut mid-word at ~200 chars, reused in OG/Twitter/subtitle VERIFIED (1 page) .
PRODUCT MOTION	Inspection.
MECHANISM	Confirms whether one template fix corrects every snippet.
OWNER	[Content/SEO] + [Dev]
DEPENDENCY	view-source (no access needed).
SUCCESS METRIC	Documented scope: % of sampled pages truncated.
WINDOW	Day 2.
FAILURE CONDITION	If only some templates truncate, scope per-template.

7D·4 Evaluate & stage the technical PRs

Days 2–7

ASSET	(a) 301 www → apex; (b) meta-description word-boundary trim.
EVIDENCE	All signals declare apex, no 301 present VERIFIED ; meta truncation VERIFIED .
PRODUCT MOTION	Evaluate → change or pass: both evaluated as needed → stage PRs. (Framework applies to every technical item.)
MECHANISM	301 consolidates signal onto apex; trim fixes every snippet at once.
OWNER	[Web/Vercel admin] + [Dev]
DEPENDENCY	Repo access + sign-off (client-gated). Staged otherwise.
SUCCESS METRIC	PRs staged/merged; www 301s; descriptions end on a word ≤155 chars.
WINDOW	Ship in 7 days if access granted; measured in 30-day.
FAILURE CONDITION	Roll back 301 if any canonical URL 4xx/5xx or rank drops >14 days.

7D·5 Configure GA4 Key Events for intent

Days 2–5

ASSET	GA4 events: <code>/pricing</code> view, <code>/demo-en</code> view, signup-start.
EVIDENCE	Key Events = 0 across the export VERIFIED .
PRODUCT MOTION	Analytics config (no site change).
MECHANISM	Makes pricing intent measurable per landing page; instrument now, analyze when volume justifies (per review #9).
OWNER	[Analytics]
DEPENDENCY	GA4 admin access (held).
SUCCESS METRIC	3 key events firing; landing-page → intent report renders.
WINDOW	Live by day 5 (forward-only).
FAILURE CONDITION	If events need a code change, fold into 7D·4 PR.

7D·6 Map cannibalization clusters + interlinking audit

Days 3–7

ASSET	~300 sitemap slugs; /industries vs /blogs overlaps; internal-link graph (orphans, money-page inlinks).
EVIDENCE	Near-duplicate clusters + vertical overlap confirmed in sitemap VERIFIED .
PRODUCT MOTION	Analysis → a merge/redirect map + an interlinking plan.
MECHANISM	One canonical page per intent; hub-and-spoke links → consolidated, well-distributed signal.
OWNER	[Content/SEO]
DEPENDENCY	GSC per-page clicks to pick cluster winners.
SUCCESS METRIC	Merge map (keep/merge/301) + interlinking plan (pillars, spokes, orphans, money-page targets).
WINDOW	Day 7 — feeds the 30-day content & interlinking work.
FAILURE CONDITION	Where two pages serve distinct intent, differentiate instead of merging.

Part Three · 30-Day plan — six workstreams

Sorted by workstream. Technical is framed as evaluate → change or pass. Owners are role placeholders for assignment.

A · TECHNICAL (EVALUATE → CHANGE OR PASS)

A1 Ship & measure 301 www → apex + meta-trim		Days 8–28
EVALUATE	Confirmed needed (no 301; meta truncates) → change.	
ASSET / EVIDENCE	Homepage split 3.9 www vs 5.5 apex VERIFIED .	
MECHANISM	301 + removal request consolidates onto apex; full-length snippets lift CTR.	
OWNER / DEP	[Web/Vercel] + [Analytics] · 7D·4 merged	
SUCCESS / WINDOW	Surviving homepage ≤3.9; dual-index resolved; stable-cohort CTR up · 14–28 days	
FAILURE	No consolidation in 28 days → audit redirect chain / canonical conflicts.	

A2 Evaluate ?fpc= params & bootstrapps subdomain		Days 10–18
EVALUATE	Are param URLs indexed / is bootstrapps a real surface? → change or pass.	
ASSET / EVIDENCE	/industries/*?fpc=; bootstrapps pos 27 INFERRED .	
MECHANISM	Self-canonical params; noindex bootstrapps if unintended.	
OWNER / DEP	[Web/Vercel] + [Dev] · confirm bootstrapps purpose; preserve ?fpc= attribution	
SUCCESS / WINDOW	Param URLs + bootstrapps drop from index · 14–21 days	
FAILURE	If intentional, retain & document (pass).	

A3 Decide on subdomain AI-visibility (docs/compliance)		Days 14–24
EVALUATE	Qualitative call (per review #2): do docs/security/compliance subdomains cost enough AI visibility to justify consolidating? → decide timing, may pass.	
ASSET / EVIDENCE	Subdomains crawled as separate properties INFERRED .	
MECHANISM	Consolidating content-bearing subdomains concentrates authority + AI-crawl surface on the root domain.	
OWNER / DEP	[Web/Vercel] + [Content/SEO] · A1 done first (apex is the priority)	
SUCCESS / WINDOW	A documented decision (consolidate / defer) with rationale · by day 24	
FAILURE	Not a hard fix — an explicit choice; default is defer until apex settles.	

A4 AEO experiments (test-measure-scale)		Days 12–30
EVALUATE	Run, don't assume (per review #8): FAQPage schema, answer-first formatting, an <code>/llms.txt</code> , crawler-token cleanup — on a test set.	
ASSET / EVIDENCE	FAQ sections without schema VERIFIED ; robots tokens dated VERIFIED .	
MECHANISM	Machine-readable Q&A + correct tokens (<code>PerplexityBot</code> , <code>ClaudeBot</code> , <code>OAI-SearchBot</code>); measure AI-referral + citation lift.	
OWNER / DEP	[Dev] + [Content/SEO] + [Analytics] · 7D·2 register; MP3 attribution	
SUCCESS / WINDOW	Test cohort vs control on AI-referral traffic; scale winners · rolling to day 30	
FAILURE	Drop approaches that don't move the metric — no sunk-cost scaling.	

B · CONTENT REVAMP — 7–9 PIECES IN MONTH ONE

B1

Revamp the priority underperformers (titles, meta, answer-first, intent)

Days 8–24, measure 14–28

#	PAGE	POS / CTR	REVAMP FOCUS
1	white-label-ai-profit-margins	4.2 / 0.04%	Title to lead with the number; surface the margin calculator
2	comparing-no-code-phone-agents-...-2026	6.7 / 0.01%	Retitle around the real latency/benchmark data inside
3	best-voice-ai-for-contact-centers	6.6 / 0.03%	Answer-first + enterprise framing
4	smith-ai-alternative-2026	6.4 / 0.03%	Lead with Smith.ai's actual pricing; merge the duplicate (B2)
5	how-to-sell-ai-chatbots-local-businesses	6.0 / 0.24%	Title + intent match; merge the "-to-" duplicate
6	white-label-ai-chatbot-pricing-comparison	6.0 / 0.11%	Pricing-first opener; real numbers (register)
7	ai-receptionist-white-label-pricing	5.2 / 0.13%	Pricing-first; internal-link to /agency
8	cheapest-ai-phone-answering-service	7.0 / 0.26%	Switch-intent angle (manual → AI), per review #7
9	best-ai-receptionist-for-small-business-2026	~7 / 0.23%	Honest named comparison + embed Trustpilot 4.6/5 (highest-impression page)
EVIDENCE	All rank ≤7 yet far below Trillet's own CTR curve VERIFIED .		
PRODUCT MOTION	Content edits, staged; publish after A1 meta-trim.		
MECHANISM	Title = searcher's phrase; answer-first, full-length snippet → CTR capture at existing rank.		
OWNER / DEP	[Content/SEO] · A1 live; 7D·2 register; 7D·1 capabilities		
SUCCESS / WINDOW	Per-page CTR toward curve (≥0.6%, stretch ≥1%) · 14–28 days post-publish		
FAILURE CONDITION	If rank drops post-edit, revert title, keep meta, re-measure.		

B2 Consolidate near-duplicate clusters (merge + 301)

Days 12–26

ASSET	Merge map (7D·6): profit-margins ×2, sell-ai-chatbots ×2, data-residency ×3, My-AI-Front-Desk ×3, Smith.ai ×2, Vapi ×3, white-label pricing cluster.
EVIDENCE	Near-duplicate slugs confirmed in sitemap VERIFIED .
PRODUCT MOTION	Merge best content into one page per intent; 301 losers.
MECHANISM	Ends self-competition; concentrates signal.
OWNER / DEP	[Content/SEO] · 7D·6 map; redirect capability
SUCCESS / WINDOW	Merged page ≥ prior best of its cluster · 14–28 days
FAILURE CONDITION	If a merge loses combined impressions after 28 days, split & differentiate.

B3 Resolve /industries vs /blogs vertical overlap

Days 16–28

ASSET	Per trade: /industries/{trade} vs /blogs/ai-answering-service-for-{trade} vs /blogs/white-label-ai-for-{trade} (plumbers, hvac, roofing first).
EVIDENCE	3 assets/trade competing; plumbers stuck at pos 22 VERIFIED .
PRODUCT MOTION	One primary intent per URL (industry = commercial; blog = informational), cross-link or merge.
MECHANISM	Stops internal competition so the strongest page reaches page 1.
OWNER / DEP	[Content/SEO] · B2 method
SUCCESS / WINDOW	/industries/plumbers pos 22 → toward page 1 · 14–28 days
FAILURE CONDITION	If demand supports both, keep both with distinct intent + links.

C · MAIN-PAGE POSITIONING & SEO STRUCTURE (2–3 PAGES)

C1 Reposition & restructure the core money pages

Days 10–28

PAGE	POSITIONING & STRUCTURE WORK
/	Lead a primary commercial keyword alongside "Voice Operating Layer"; clear SMB / Agency / Enterprise routing above the fold; Organization + Product schema; internal links to money pages
/pricing	Consolidate with /plans (one canonical); Product/Offer schema; single source-of-truth numbers (register); FAQ block + schema
/agency	Sharpen the agency positioning (best-converting segment — white-label long-tail already at 0.39% CTR); structure for "white label / agency" intent; link from high-authority blogs
EVIDENCE	Homepage title carries no commercial term VERIFIED ; /pricing & /plans both indexed VERIFIED ; agency/white-label is the best-converting cluster VERIFIED .
PRODUCT MOTION	Positioning + on-page structure + schema; content via register.
MECHANISM	Commercial-keyword positioning + segment clarity + schema improves relevance, CTR, and conversion routing.
OWNER / DEP	[Content/SEO] + [Product Marketing] · 7D·1, 7D·2
SUCCESS / WINDOW	Homepage ranks & lifts CTR for a commercial head term; /pricing single canonical; /agency engaged-session rate up · 21–35 days
FAILURE CONDITION	If a positioning change dips brand-term CTR, revert the hero, keep the structural/schema gains.

D · INTERLINKING

D1 Build hub-and-spoke + money-page internal linking		Days 14–30
ASSET	Pillars (cost hub, compliance pillar, comparison hub, whitelabel-guide) ↔ their spoke posts; high-authority pages (homepage, call-forwarding-australia, best-ai-receptionist) → money pages (/pricing, /agency, /demo-en); /industries ↔ related blogs; orphan pages.	
EVIDENCE	~300 posts with weak internal structure; money pages under-linked; industry/blog overlap VERIFIED .	
PRODUCT MOTION	Contextual internal links per the 7D·6 interlinking plan; fix orphans.	
MECHANISM	Distributes authority to money pages + pillars; helps the strongest page in each cluster rank; improves crawl paths.	
OWNER / DEP	[Content/SEO] · 7D·6 plan; B2/B3 consolidation first (link to survivors)	
SUCCESS / WINDOW	Every money page + pillar has ≥5 relevant internal links; 0 orphans in the priority set · 14–28 days	
FAILURE CONDITION	Avoid over-linking / exact-match anchor stuffing — keep anchors natural.	

E · NEW CONTENT (SELECTIVE — CONSOLIDATION-FIRST)

E1 Cost/pricing hub · Voicify + comparison hub · compliance pillar		Days 12–35
ASSET	(1) canonical "AI answering service cost" hub (absorbs scattered pricing posts; 195-query cluster); (2) <code>trillet-vs-voicify</code> + a <code>/compare</code> hub over the consolidated competitor pages; (3) one compliance/data-residency pillar (merges the trio; leads with "included on every plan").	
EVIDENCE	Cost cluster 49,503 impr / 0.09% CTR; compliance cluster 5,885 impr / 0.00% CTR; Voicify demand with no page VERIFIED / INFERRED .	
PRODUCT MOTION	New/merged pillars with FAQ schema; 301 thin posts in; answer-first.	
MECHANISM	Captures underserved commercial + compliance intent Trillet already ranks for; pillars anchor the interlinking (D1).	
OWNER / DEP	[Content/SEO] (+ [Legal] for compliance) · 7D·2 register; B2 for competitor consolidation	
SUCCESS / WINDOW	Cost hub page 1 for "ai answering service cost"; compliance cluster clicks >0; Voicify indexed · 21–35 days	
FAILURE CONDITION	No compliance claim ships without legal sign-off — hard gate. If cost hub cannibalizes /pricing, canonicalize appropriately.	

F · SAAS-MARKETPLACE PROFILES & OFF-PAGE

F1 Create/complete profiles · seed reviews · attribute

Days 8–30

ASSET	G2, Capterra + GetApp + Software Advice (one Gartner submission), TrustRadius, Product Hunt; review-request flow; UTM'd profile links + GA4 AI-referral channel.
EVIDENCE	No presence today VERIFIED ; Trustpilot 4.6/5 (20) exists VERIFIED ; competitors listed INFERRED .
PRODUCT MOTION	Claim/complete to 100%; ToS-compliant review seeding from happy customers; tag + measure referral + AI-referral.
MECHANISM	Trust signals, backlinks, review rich-snippets, and the sources AI answer engines cite for "best AI receptionist / voice AI".
OWNER / DEP	[Product Marketing] + [Analytics] · 7D·2 register; 7D·5 events
SUCCESS / WINDOW	5 profiles 100% complete (day 20); ≥10 reviews each on G2 + Capterra (day 30); referral + AI channel visible
FAILURE CONDITION	Never incentivize reviews against platform ToS — delisting/reputation risk.

Part Four • Trackers

Two simple views, updated in Sheets: a **task tracker** (run the plan — owner, dependency, status, window, success metric per action) and a **metrics tracker** (update the data — baseline, target, periodic readings, linked task). Every material number carries source, date range, filters and supporting artifact; every conclusion labeled verified / inferred / unverified.

Seed baselines (GSC, Web): non-brand share 3.5% · striking-distance non-brand impressions 155,009 · stable-cohort CTR 3.74% · homepage www 3.9 / apex 5.5 · /industries/plumbers 22.3 · cost cluster 49,503 impr @ 0.09%. Trust: Trustpilot 4.6/5 (20). GA4 (Jun 15–Jul 12): 562 users, key events 0 (pre-config baseline).