

CONSOLIDATED FINDINGS & ACTION PLAN

Trillet.ai — SEO, Content & Keyword Master Report

The major, non-trivial findings from three companion audits — technical/AEO/GEO, twelve months of Search Console performance, and a content-and-keyword revamp brief — brought together into one prioritized plan.

Compiled – July 8, 2026

Scope – trillet.ai (7 hostnames), 305 blog posts, 31 industry pages, 12 months of GSC data

Method – verified DNS/hosting data, the site's own Search Console export, and Google's indexed understanding of page content (direct crawling was blocked throughout by this session's network policy)

Source 1 — SEO, AEO & GEO Audit (search-intelligence + DNS verification)

Source 2 — Content & Keyword Performance Analysis (12-month GSC export)

Source 3 — Content Revamp Brief (page-level and keyword-cluster fixes)

Executive summary

Twelve months of real data, not projection. The site ranks — the problem is almost entirely architectural and editorial, not a ranking failure.

CLICKS / 12MO 14,100	IMPRESSIONS 1.88M	SITE-WIDE CTR 0.75%	BRAND SHARE OF CLICKS 93.3%
--------------------------------	-----------------------------	-------------------------------	--

Three findings drive everything else in this report:

1. **Ranking signal is split across up to seven hostnames.** Half of all indexed URLs (216 of 432) exist under more than one host, holding **89.6% of all impressions**. This is the single highest-leverage fix available and requires no content rewriting.
2. **A January 2026 content launch inflated impressions 8–40× while CTR collapsed under 1% and stayed there.** Position didn't get worse — this is a title/snippet relevance failure at scale, not a ranking failure.
3. **The content itself is frequently good and the packaging is what's failing.** Real differentiators (compliance included free, competitive pricing, benchmarked latency data) exist in the copy and simply aren't reaching the exact phrases buyers search.

02 – HIGHEST LEVERAGE

Site architecture: the canonicalization problem, proven

Flagged first from indexation patterns, then confirmed directly with real click/impression/position data.

PATHS DUPLICATED ACROSS HOSTS 216 / 432	IMPRESSIONS ON SPLIT PATHS 89.6%	DISTINCT HOSTNAMES FOUND 7	HOMEPAGE POSITION GAP, APEX VS WWW 1.6
--	---	---	---

Seven hostnames splitting one property

HIGH

trillet.ai, www.trillet.ai, docs.trillet.ai, app.trillet.ai, security.trillet.ai, certification.trillet.ai, and bootstrapps.trillet.ai (ranking at position 27 for a homepage query — reads like an unintended exposure rather than a deliberate property, worth checking whether it should be indexed at all). DNS-verified: the whole property runs on Vercel, so host consolidation is a dashboard setting, not custom redirect code.

The homepage split, concretely

HIGH

www.trillet.ai/ — 7,863 clicks / 54,438 impr. / position 3.9. trillet.ai/ (apex) — 1,094 clicks / 7,448 impr. / position 5.5. Identical content, 1.6 positions apart — direct proof the two hosts are splitting, not sharing, ranking signal.

The same split repeats across the entire blog

HIGH

E.g. `best-ai-receptionist-for-small-business-2026` : 117,316 combined impressions, `www` at position 6.9 vs. `apex` at 8.2. `cheapest-ai-phone-answering-service` : `www` at 7.0 vs. `apex` at 14.6 — more than a full page apart. 216 such pairs exist, most of them blog posts.

Highest-demand industry page stuck on page 3

SERIOUS

`/industries/plumbers` — the highest-impression-volume industry page (9,095) — sits at an average position of 22.3, split 19.8 (`www`) vs. 27.9 (`apex`). This should realistically compete for page 1 given the demand it already attracts.

Tracking parameters getting indexed as separate URLs

SERIOUS

Two `/industries/*` pages are indexed with a `?fpc=` referral-tracking token intact — without a canonical tag, Google can index unbounded parameter permutations of the same page.

Duplicate pricing pages

MEDIUM

Both `/pricing` and `/plans` are indexed, competing for the same commercial query instead of consolidating authority. Real fragmentation compounds this: figures of \$29, \$49, \$99, and \$299/month appear across different pages and blog mentions.

Rendering risk for AI crawlers

MEDIUM – UNVERIFIED

If the Vercel-hosted site renders content client-side, most AI answer-engine crawlers (GPTBot, ClaudeBot, PerplexityBot) — which mostly don't execute JavaScript — could see a near-empty page even while Googlebot indexes it fine. The single highest-leverage open question for GEO; a plain view-source check resolves it.

03

The 2026 content flood, and where clicks actually come from

PERIOD	IMPRESSIONS/MO	CTR
Jul–Nov 2025 (baseline)	5,300 – 8,100	9 – 16%
Dec 2025 (hinge month)	9,828	4.1%
Jan–Jul 2026 (post-launch)	35,700 – 337,900	0.6 – 1.0%

Ranking position didn't get meaningfully worse in this window (often better). This is a relevance/snippet failure at scale, consistent with the mass rollout of the 305-post blog and 31-page industry cluster (most titled "2026") landing almost entirely in this window.

SEGMENT	% OF CLICKS	% OF IMPRESSIONS	CTR
Homepage	65.0%	3.7%	13.3%
Blog (305 posts)	20.6%	84.2%	0.18%
Industry pages (31)	0.9%	2.1%	0.31%

Two-thirds of every click comes from people typing the brand name. The blog carries 84% of impression volume and returns 21% of clicks — the content is doing the hard part (getting found) and failing at the cheap part (getting clicked).

Worst-performing pieces (real posts, ranking page 1, earning almost nothing)

POST	IMPR.	POSITION	CTR
comparing-no-code-phone-agents-...-2026	8,819	6.7	0.01% (1 click)
ai-answering-service-guide	25,227	30.7	0.03%
voice-ai-data-residency-requirements-by-region	34,864	6.5	0.04%
voice-ai-white-label-pricing-breakdown-2026	46,312	6.2	0.06%

04

Content & keyword revamp: what the real search terms show

Every cluster below shows the same shape: decent-to-good rankings, real search demand, near-zero clicks. Not 305 separate content problems — one problem (title/snippet mismatch) showing up 305 times.

CLUSTER	IMPR.	CTR	READ
Cost / cheap / pricing	49,503	0.09%	BIGGEST GAP
Compliance / data / security	5,885	0.00%	REAL STRENGTH, ZERO CONVERSION
Industry-specific	34,712	0.07%	CONFIRMS PAGE-LEVEL FINDING
Competitor pricing	1,688	0.00%	INTENT NOT ANSWERED FIRST
White-label / agency	16,019	0.39%	BEST OF THE SET – PROOF THE FIX WORKS
Generic head terms	10,258	~0.01%	NOT WINNABLE YET – DEPRIORITIZE

Confirmed keyword cannibalization

At least two competitors are covered by multiple, separately-ranking posts instead of one consolidated page: **Smith.ai** (`smith-ai-alternative-2026` and `trillet-vs-smith-ai-comparison-2026`) and **Vapi** (three separate posts). All 27 comparison/alternative posts should be audited for the same pattern.

The revamp checklist

1. Open with the answer, not the pitch — the first sentence should satisfy the exact query.
2. Title = the searcher's phrase, not brand voice.
3. One piece of content per search intent — merge cannibalized pairs.
4. Match specificity to competitiveness — go deeper on long-tail terms already winning; stop pouring effort into unwinnable head terms.
5. State the actual number (price, certification) in the first screen — never "contact us."
6. Add FAQPage/HowTo schema wherever content already answers a question.
7. Link cost/compliance/comparison content to `/pricing` and the compliance page.

8. Consolidate before expanding — fix existing pages before adding new ones.

Highest-ROI single fix: rewrite titles/meta on the worst-CTR posts — they collectively hold over 160,000 impressions at page-1 positions and return almost nothing. Estimated floor opportunity: **4,000+ additional clicks** over a comparable 12-month period from title/snippet fixes alone, before any new content is written (conservative 2% CTR benchmark vs. the 0.27% currently captured on non-brand queries; real long-tail volume is larger than what GSC's 1,000-row export cap captured, so this is a floor, not a ceiling).

05

AEO & GEO readiness

Whether Trillet shows up in Google's AI Overviews and featured snippets (AEO), and whether ChatGPT, Perplexity, Gemini, and Claude actually cite it (GEO).

- **Already working:** question-shaped blog titles, real comparison tables (the right shape for snippet extraction), and rich results already appearing in Search Console (Review snippet, Product snippet impressions confirmed) — proof structured data works somewhere on the site already.
- **Unverified, needs direct access:** FAQPage/HowTo schema coverage, actual `robots.txt` AI-crawler directives (GPTBot, ClaudeBot, PerplexityBot, Google-Extended).
- **Missing:** no `/llms.txt` — a low-cost, curated summary file for AI crawlers.
- **The open technical question:** client-side rendering risk (\$2) — could mean strong Google indexing and near-zero AI-crawler visibility simultaneously.
- **Third-party citation gap:** zero Reddit, Quora, or Hacker News discussion found — AI answer engines weight forum consensus heavily. **Genuine strength:** independent, third-party YouTube creator videos exist, not just brand-published content — worth actively encouraging more of, since it's already working.

06

Off-page & trust signals

Correction to the initial audit: Trillet does have a Trustpilot listing — **4.6/5 from 20 reviews** — confirmed on a later pass. Real, if thin, social proof that isn't surfaced anywhere in the site's own content.

- **Still missing:** G2, Capterra, and Product Hunt listings — competitors (Retell AI, Vapi, Bland AI, Synthflow) are actively reviewed on these; a real gap in trust signals and backlink sourcing.
- **Existing footprint:** LinkedIn company page and a GitHub organization (`github.com/TrilletAI`) — confirm both are linked from the site footer for entity consistency.
- **Brand-name collision:** a third-party directory lists an implausible "\$388M raised" figure for Trillet, almost certainly conflated with similarly-named companies ("Rillet," an AI ERP startup; "Triller," the social app). Worth requesting a correction — bad aggregator data can leak into AI answer engines.

07

From a first-time visitor's perspective

- **The tagline may not orient a first-time visitor.** "The Voice Operating Layer" is memorable but abstract — it doesn't signal which of three very different offers (SMB, agency, enterprise) applies to the visitor, compounded by a homepage title with no commercial keyword.
- **Fragmented pricing is a trust risk, not just a technical one.** \$29/\$49/\$99/\$299 figures scattered across pages erode trust even if each is a genuinely different tier — one canonical pricing page fixes both the technical and the visitor-facing problem at once.
- **Programmatic industry pages risk reading as doorway pages.** A visitor landing on a templated page with no vertical-specific substance bounces faster than one with even one real example or testimonial from that trade — the same fix (unique per-vertical substance) solves the SEO thin-content risk and the visitor-trust risk simultaneously.

08 — THE PLAN

Consolidated action plan

Do first — structural, affects everything downstream

- 1 Consolidate all seven hostnames to one canonical domain via Vercel's domain settings — 301s, self-referencing canonicals, sitemap update. Highest-leverage fix in the entire body of work.
- 2 Confirm whether AI crawlers can read rendered content (client-side vs. server-rendered check).
- 3 Fix `?fpc=` tracking-parameter indexation on `/industries/*` pages.
- 4 Consolidate `/pricing` and `/plans` into one canonical URL and number.

Do next — content, high ROI

- 1 Rewrite titles/meta on the worst-CTR posts first (§3) — hours of work, ~160,000 impressions currently wasted.
- 2 Rebuild the cost/pricing cluster to state real numbers plainly and match "cheap/low-cost" search phrasing directly.
- 3 Rebuild the compliance cluster to lead with the HIPAA/TCPA/ACMA/GDPR-included claim in the title and first sentence.
- 4 Merge cannibalized comparison posts (Smith.ai, Vapi confirmed; audit all 27 for the same pattern).
- 5 Add FAQPage/HowTo schema; publish an explicit AI-crawler `robots.txt` policy; add `/llms.txt` .
- 6 Rewrite the homepage title/meta to lead with a primary commercial keyword.

Ongoing

- 1 Get listed and collecting reviews on G2, Capterra, and Product Hunt; surface the real Trustpilot 4.6/5 rating directly in comparison and "best of" content.
- 2 Seed genuine third-party discussion (Reddit, more creator videos) — the YouTube presence already proves this works.
- 3 Hold off on `/industries/*` or comparison pages until the existing 31 industry pages and 27 comparison posts are fixed.
- 4 Set an annual refresh cadence for year-stamped ("2026") content.
- 5 Request a correction from the third-party directory on the incorrect funding figure.